

Exploring Binary

What Powers of Two Look Like Inside a Computer

By [Rick Regan](#) January 7th, 2009

A [power of two](#), when expressed as a binary number, is easy to spot: it has one, and only one, 1 bit. For example, 1000, 10, and 0.001 are powers of two. Inside a computer, however, numbers are more generally represented in binary code, not as “pure” binary numbers. As a result, you may not be able to look at the binary representation of a number and tell at a glance whether it’s a power of two or not; it depends on how it’s encoded.

Numbers in computers are encoded in several forms, the two most common being *binary integer* and *floating-point binary*. Powers of two are represented no differently, but because of their special relationship to binary notation, have a distinct look.

We’ll start by reviewing the binary integer and floating-point binary representations, and then we’ll see how powers of two are encoded in those forms.

Binary Integers

The standard way to represent integers is with [two’s complement](#) notation. The two’s complement representation of a nonnegative integer is faithful to its representation as a pure binary number, with some superfluous differences; it will not have commas or spaces and may have leading zeros added.

The two’s complement representation of a negative integer is *not* faithful to its representation as a pure binary number. It has the equivalent of a minus sign, which is a 1 bit in the most significant or leftmost position, but the value in the rest of the field is [not the pure binary representation of the integer](#). (That’s all I’ll say about it in this article, since negative numbers are not powers of two.)

Integers come in several sizes: 8 bits, 16 bits, 32 bits, and 64 bits are the common choices. The size determines the number of integers that can be represented. For example, a 32-bit integer can represent 2^{32} values.

An integer can be unsigned or signed. Unsigned integers represent only nonnegative values. A 32-bit unsigned integer can represent an integer between 0 and $2^{32} - 1$. Signed integers represent both negative and nonnegative values, with an equal number of each representable. A 32-bit signed integer can represent an integer between -2^{31} and $2^{31} - 1$.

Floating-Point Binary Numbers

The standard way to represent real numbers is with [IEEE 754 floating-point binary](#). Floating-point representation is akin to scientific notation; it separates a number into a sign part, an exponent part, and a fraction part. This means that a floating point number is not stored as a pure binary number.

The two common forms of floating-point are single-precision, which is 32 bits, and double-precision, which is 64 bits. The forms differ by the size of the exponent and fraction fields; they are bigger in double-precision.

Single-Precision Floating-Point

Single-precision floating-point numbers are 32 bits long, and consist of three fields:

1. **Sign.** The sign field is 1 bit. It is 0 for positive numbers, 1 for negative numbers, and can be 0 or 1 for the number 0.
2. **Exponent.** The exponent field is 8 bits. It represents the exponent n to which 2 is raised, scaling the fraction by 2^n . The exponent is stored in biased form, which is a special encoding that stores positive and negative exponents as positive integers. To compute the exponent, interpret this field as an unsigned binary integer and subtract 127, giving a signed binary integer result. For example, a value of 10000011 (131) represents an exponent of 4, and a value of 01111101 (125) represents an exponent of -2. The values 0 and 255 have special meanings, and are not interpreted as exponents.
3. **Fraction.** The fraction field is 23 bits. If the number is normalized, the exponent field will be nonzero, and there will be an implicit 1 bit preceding the fraction. For example, 2^2 has an exponent field value of 10000001 (129), which means the exponent is 2. The number is interpreted as 1.0×2^2 , but only the 0 is stored in the fraction.

If the number is denormalized (AKA denormal or subnormal), the exponent field will be zero, and there will *not* be an implicit 1 bit preceding the fraction. A denormalized number has an implicit starting exponent of -126, minus the position of the first 1 bit in the fraction field. For example, 2^{-128} has an exponent field value of zero, and a fraction field value of .01. The number is interpreted as 0.01×2^{-126} , which is 2^{-128} .

Double-Precision Floating-Point

Double-precision floating-point numbers are 64 bits long, but are otherwise similar to single-precision numbers. The only differences are in the exponent and fraction fields. The exponent field is 11 bits, has a bias of 1023, and gives the values 0 and 2047 special meanings. The fraction field is 52 bits.

Normalized and denormalized numbers are represented similarly as in single-precision floating-point. For example, 2^2 , which is a normalized number, has an exponent value of 10000000001 (1025) and a fraction field value of 0. The number 2^{-1024} , which is a denormalized number, has an exponent value of zero and a fraction field value of .01.

Powers of Two in Integers

Powers of Two in Unsigned Integers

An n-bit unsigned integer can represent n powers of two; specifically, the nonnegative powers of two from 2^0 through 2^{n-1} . For example, a 32-bit unsigned integer can represent the 32 powers of two from 2^0 through 2^{31} :

Powers of Two in a 32-bit Unsigned Integer

Power of Two	32-Bit Unsigned Integer Representation
2^0	00000000000000000000000000000001
2^1	00000000000000000000000000000010
2^2	00000000000000000000000000000100
...	...
2^{29}	00100000000000000000000000000000

2^{30}	01000000000000000000000000000000
2^{31}	10000000000000000000000000000000

From the table you can deduce that an unsigned integer is a power of two if and only if it has exactly one 1 bit.

Powers of Two in Signed Integers

An n -bit signed integer can represent $n-1$ powers of two; specifically, the nonnegative powers of two from 2^0 through 2^{n-2} . For example, a 32-bit signed integer can represent the 31 powers of two from 2^0 through 2^{30} :

Powers of Two in a 32-bit Signed Integer

Power of Two	32-Bit Signed Integer Representation
2^0	00000000000000000000000000000001
2^1	00000000000000000000000000000010
2^2	00000000000000000000000000000100
...	...
2^{29}	00100000000000000000000000000000
2^{30}	01000000000000000000000000000000

The value 10000000000000000000000000000000, which is 2^{31} in an unsigned integer, is -2^{31} in a signed integer; thus, it's not a power of two in a signed integer.

From the table you can deduce that a signed integer is a power of two if and only if its sign bit is 0 and the remaining part of its binary representation has exactly one 1 bit.

Powers of Two in Floating-Point

Powers of Two in Single-Precision

Single-precision floating-point can represent 277 powers of two, from 2^{-149} through 2^{127} . Denormalized powers of two are those from 2^{-149} through 2^{-127} ; normalized powers of two are those from 2^{-126} through 2^{127} . This table shows the range of powers of two covered and their encodings:

Powers of Two in a 32-Bit Floating-Point Number

Power of Two	32-Bit Floating-Point Representation		
	–	Exponent	Fraction
2^{-149}	0	00000000	000000000000000000000001
2^{-148}	0	00000000	000000000000000000000010
2^{-147}	0	00000000	000000000000000000000100
...
2^{-128}	0	00000000	010000000000000000000000
2^{-127}	0	00000000	100000000000000000000000
2^{-126}	0	00000001	000000000000000000000000
2^{-125}	0	00000010	000000000000000000000000
2^{-124}	0	00000011	000000000000000000000000
...
2^{-2}	0	01111101	000000000000000000000000
2^{-1}	0	01111110	000000000000000000000000
2^0	0	01111111	000000000000000000000000
2^1	0	10000000	000000000000000000000000
2^2	0	10000001	000000000000000000000000
2^3	0	10000010	000000000000000000000000
2^4	0	10000011	000000000000000000000000
...
2^{125}	0	11111100	000000000000000000000000
2^{126}	0	11111101	000000000000000000000000
2^{127}	0	11111110	000000000000000000000000

From the table you can deduce that a single-precision floating-point number is a power of two if and only if

Try it Out

I wrote a C program to [print the fields of a double-precision floating-point-number](#). You can use that program to explore the representation of powers of two further.



EB

Related

- [Displaying IEEE Doubles in Binary Scientific Notation](#)
- [Displaying the Raw Fields of a Floating-Point Number](#)
- [A Simple C Program That Prints 2,098 Powers of Two](#)
- [Ten Ways to Check if an Integer Is a Power Of Two in C](#)
- [How to Check If a Number Is a Power of Two](#)



Binary Subtraction	The Powers of Two	Number of Bits in a Decimal...	Binary Addition
Floating-Point Will Still Be Broken In...	What a Binary Counter Looks and Sounds...	My Fascination with Binary Numbers	I I I

-  [Get articles by RSS \(What Is RSS?\)](#)
-  [Get articles by e-mail](#)

3 comments

1. **krupal**

September 8, 2014 at 11:25 am

Thanks a lot...This solved my all issues in understanding floating point numbers.

2. **AndyF**

April 8, 2017 at 7:02 pm

Hi, there's a small typo in this sentence from your above article: "A 32-bit signed integer can represent an integer between -2^{31} and $2^{31} - 1$."

It should be – (minus) 2 raised to the power of positive 31, not to the power of negative 31. (Apologies, I can't reproduce the superscript in this comment, so I used the '^' character to indicate "to the power of".)

3. **Rick Regan**

April 8, 2017 at 10:15 pm

@AndyF,

Thanks for spotting that — I just fixed it. It was there for 8 years!

Comments are closed.

Copyright © 2008–2020 Exploring Binary

[Privacy policy](#)

Powered by [WordPress](#)

